

SNPsea Reference Manual

Kamil Slowikowski

December 31, 2013

Contents

Introduction	2
Contact	2
Visual Summary	3
Installation	3
Data	4
C++ Libraries	5
Python Packages	6
R Packages	7
Usage	7
Example	7
Options	8
Input File Formats	9
Output Files	12
Output Visualizations	15

Introduction

SNPsea is a nonparametric permutation analysis for identifying pathways and tissue types influenced by the alleles discovered through genome-wide association studies (GWAS). It was originally conceived to test for enrichment of cell type-specific expression of genes in loci near trait-associated SNPs.

The implementation described here is generalized, so you may provide a quantitative gene matrix with gene expression (or any other measurements) or a binary gene matrix with presence absence (1, 0) values. The columns of the matrix might be tissues, cell types, GO annotation codes, or any other types of conditions. In general, this analysis is appropriate when you are interested in testing for enrichment of condition-specificity of genes linked to a given set of trait-associated SNPs.

The following hypothesis is tested by this analysis:

If trait-associated alleles impact a small number of pathogenic tissues or cell types, then the subset of genes with critical functions in those pathogenic cell types are likely to be within trait-associated loci.

We assume that a gene's specificity to a given cell type or condition is a reasonable indicator of the gene's importance to the function of that cell type.

Please see the following publications for additional information outside the scope of this reference manual:

Slowikowski, K. et al. SNPsea: test trait-associated loci for enrichment of condition-specificity of gene measurements or binary annotations. Manuscript in progress.

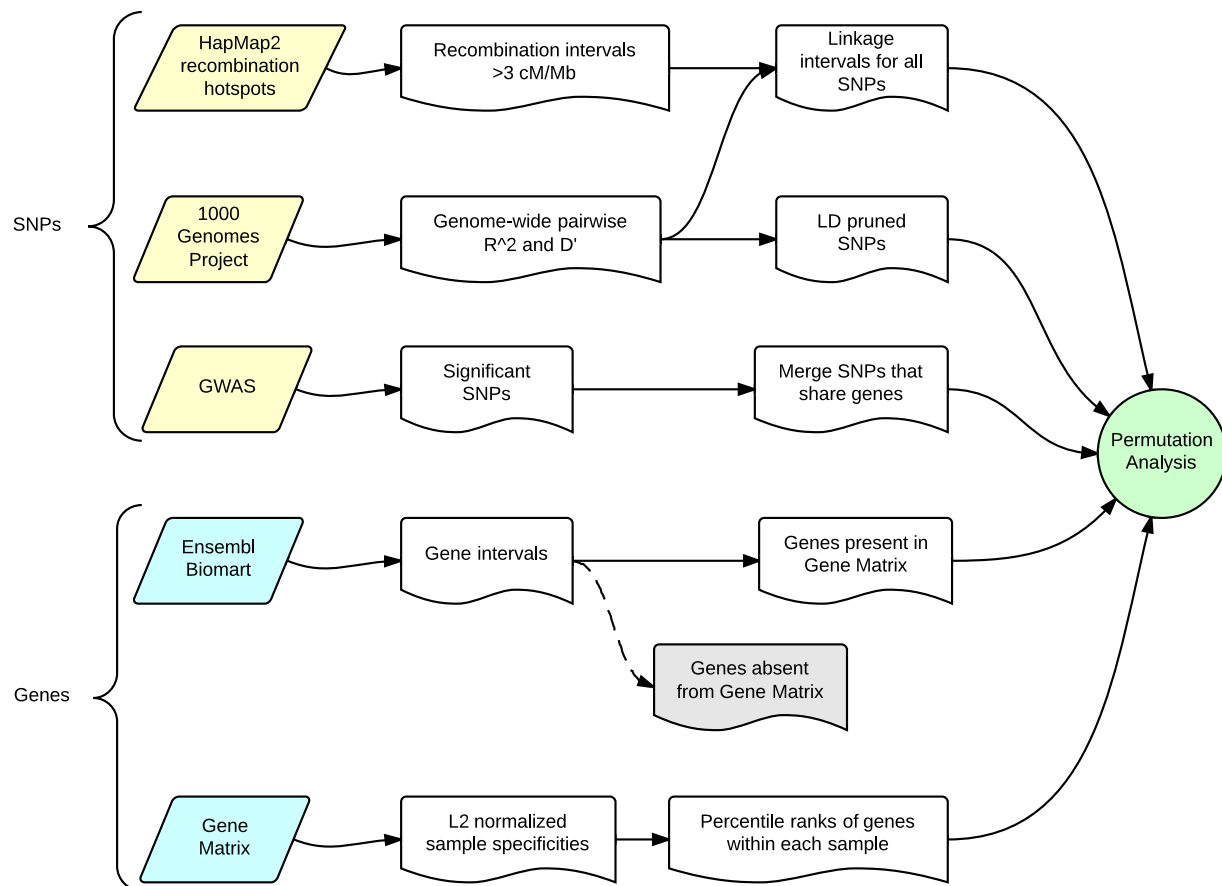
Hu, X. et al. Integrating autoimmune risk loci with gene-expression data identifies specific pathogenic immune cell subsets. *The American Journal of Human Genetics* 89, 496–506 (2011). [PubMed](#)

Contact

Please contact me with questions and comments: slowikow@broadinstitute.org

Visual Summary

Flow Chart



This flow chart shows the input data required to perform the analysis, and a summary of the intermediate steps.

Installation

Linux 64-bit

Download the binary and run it: <https://github.com/slowkow/snpsea/releases>

All other platforms

Download the source code: <https://github.com/slowkow/snpsea>

After unzipping the source code, run:

```
# Ubuntu
sudo apt-get install build-essential libopenmpi-dev libgsl0-dev
```

```

# Mac
# Install MacPorts: http://www.macports.org
# Then do this:
sudo port selfupdate
sudo port install gcc48 openmpi gsl

# Broad Institute
# You can add this to your .my.bashrc or .my.cshrc
use .gcc-4.8.1 .openmpi-1.4 .gsl-1.14

# Compile the code.
cd snpsea/src
make

# Move the generated executable wherever you like.
mv ../bin/snpsea ~/bin/

```

Data

Download the compressed archive with data required to perform this analysis here (138M):

<http://dx.doi.org/10.6084/m9.figshare.871430>

Contents:

```

G02013.gct.gz
ImmGen2012.gct.gz
LDL_Teslovich2010.txt
Lango2010.txt.gz
NCBIGenes2013.bed.gz
NovartisGeneAtlas2004.gct.gz
TGP2011.bed.gz

```

G02013.gct.gz

A GCT formatted gene matrix with 1s and 0s indicating presence or absence of 19,111 genes in 1,751 Gene Ontology annotations.

ImmGen2012.gct.gz

Gene expression data for 15,139 genes across 249 blood cell types from [GSE15907](#). Replicates for each cell type were averaged. For each gene, the single probe with the largest minimum was selected.

Immunological Genome Project. <http://www.immgen.org/>

LDL_Teslovich2010.txt

37 SNPs taken from:

Teslovich TM, Musunuru K, Smith AV, et al. Biological, clinical and population relevance of 95 loci for blood lipids. Nature. 2010;466(7307):707-13. [PubMed](#)

Lango2010.txt.gz

A list of 56,890 SNPs pruned by linkage disequilibrium that span the whole genome. Null SNP sets matched on the number of genes in the user's SNP set are sampled from this list. See this paper for more information:

Lango allen H, Estrada K, Lettre G, et al. Hundreds of variants clustered in genomic loci and biological pathways affect human height. Nature. 2010;467(7317):832-8. [PubMed](#)

NCBIgenes2013.bed.gz

40,437 Gene coordinates were obtained from the [NCBI](#). The coordinates I provide are a subset of those listed in [gene2refseq.gz](#).

NovartisGeneAtlas2004.gct.gz

Gene expression data for 17,581 genes across 79 human tissues from [GSE1133](#). Replicates for each tissue were averaged. For each gene, the single probe with the largest minimum was selected.

Su AI et al. A gene atlas of the mouse and human protein-encoding transcriptomes. Proc Natl Acad Sci U S A, 2004 Apr 9;101(16):6062-7

TGP2011.bed.gz

Linkage intervals for a filtered set of 22,518,294 SNPs from the [1000 Genomes Project](#) Phase 1 (May 21, 2011). SNP genotypes were obtained from the [BEAGLE](#) release v3 website and processed to create linkage intervals for each SNP. The linkage intervals were extended to the nearest [HapMap](#) recombination hotspot with >3 cM/Mb recombination rate.

C++ Libraries

To compile SNPsea, you will need a modern C++ compiler that supports [c++0x](#) and the libraries listed below.

See: [Installation](#)

[Eigen](#)

Eigen is a C++ template library for linear algebra: matrices, vectors, numerical solvers, and related algorithms.

OpenMPI

MPI is a standardized API typically used for parallel and/or distributed computing. Open MPI is an open source, freely available implementation.

GSL - GNU Scientific Library

The GNU Scientific Library (GSL) is a numerical library for C and C++ programmers.

GCC, the GNU Compiler

The GNU Compiler Collection is a compiler system produced by the GNU Project supporting various programming languages.

I use `c++0x` features in my C++ code, so you must use a compiler that supports them. I compiled successfully with versions 4.6.3 (the default version for Ubuntu 12.04) and 4.8.1.

Python Packages

To plot visualizations of the results, you will need Python 2.7 and the packages listed below.

Instructions: Install with `pip`:

```
pip install docopt numpy pandas matplotlib
```

Note: The packages available on the Ubuntu repositories may be outdated and might fail to work. So, avoid using `apt-get` for these dependencies.

docopt

Command-line interface description language.

numpy

NumPy is the fundamental package for scientific computing with Python.

pandas

pandas is an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the Python programming language.

matplotlib

matplotlib is a python 2D plotting library which produces publication quality figures in a variety of hardcopy formats and interactive environments across platforms.

Note: On a server with no display, please edit your [matplotlibrc](#) file to use the **Agg** backend. Otherwise, you may see an error message like this:

```
_tkinter.TclError: no display name and no $DISPLAY environment variable
```

R Packages

Some visualizations use R and ggplot2 instead of Python and matplotlib.

Instructions: Start a session in R and run:

```
install.packages(c("data.table", "reshape2", "gap", "ggplot2"))
```

[data.table](#)

Extension of data.frame for fast indexing, fast ordered joins, fast assignment, fast grouping and list columns.

[reshape2](#)

Flexibly reshape data: a reboot of the reshape package.

[gap](#)

Genetic analysis package.

[ggplot2](#)

An implementation of the Grammar of Graphics.

Usage

Example

Here is a [Bash](#) script with a usage example:

```

options=(
  --snps LDL_Teslovich2010.txt
  --gene-matrix NovartisGeneAtlas2004.gct.gz
  --gene-intervals NCBIgenes2013.bed.gz
  --snp-intervals TGP2011.bed.gz
  --null-snps Lango2010.txt.gz
  --out out
  --slop 250e3
  --threads 4
  --null-snpsets 1e3
  --min-observations 50
  --max-iterations 1e6
)
snpsea ${options[*]} &> log.txt

```

This will run the analysis on SNPs associated with LDL cholesterol and test for tissue-specific expression of the nearby genes across 79 human tissues in the Novartis 2011 gene expression matrix. Additionally, 1000 null random matched SNP sets will be tested and their results will also be recorded. Each tissue will be tested up to 1 million times, or testing will stop for a tissue if 50 matched SNP sets are observed to achieve a higher specificity score than the user's SNPs.

Options

All input files may be optionally compressed with [gzip](#).

Required

<code>--snps ARG</code>	Text file with SNP identifiers in the first column. Instead of a file name, you may use 'randomN' with an integer N for a random SNP list of length N.
<code>--gene-matrix ARG</code>	Gene matrix file in GCT format. The Name column must contain the same gene identifiers as in <code>--gene-intervals</code> .
<code>--gene-intervals ARG</code>	BED file with gene intervals. The fourth column must contain the same gene identifiers as in <code>--gene-matrix</code> .
<code>--snp-intervals ARG</code>	BED file with all known SNP intervals. The fourth column must contain the same SNP identifiers as in <code>--snps</code> and <code>--null-snps</code> .
<code>--null-snps ARG</code>	Text file with names of SNPs to sample when

generating null matched or random SNP sets.
These SNPs must be a subset of --snp-intervals.

--out ARG Create output files in this directory. It will be created if it does not already exist.

Optional

--condition ARG Text file with a list of columns in --gene-matrix to condition on before calculating p-values. Each column in --gene-matrix is projected onto each column listed in this file and its projection is subtracted.

--slop ARG If a SNP interval overlaps no gene intervals, extend the SNP interval this many nucleotides further and try again.
[default: 250000]

--threads ARG Number of threads to use.
[default: 1]

--null-snpsets ARG Test this many null matched SNP sets, so you can compare your results to a distribution of null results.
[default: 10]

--min-observations ARG Stop testing a column in --gene-matrix after observing this many null SNP sets with specificity scores greater or equal to those obtained with the SNP set in --snps. Increase this value to obtain more accurate p-values.
[default: 25]

--max-iterations ARG Maximum number of null SNP sets tested for each column in --gene-matrix. Increase this value to resolve smaller p-values.
[default: 1000]

Input File Formats

--snps ARG

You must provide one or more comma-separated text files. SNP identifiers must be listed one per line. Only the first column is used.

head LDL_Teslovich2010.txt

```
rs11136341 chr8 145043543
rs3757354 chr6 16127407
rs12027135 chr1 25775733
rs217386 chr7 44600695
rs1169288 chr12 121416650
rs7225700 chr17 45391804
rs2479409 chr1 55504650
rs247616 chr16 56989590
rs2954022 chr8 126482621
rs1564348 chr6 160578860
```

Instead of providing a file with SNPs, you may use “randomN” like this:

```
--snps random20
```

to sample 20 random SNPs from the **--snp-intervals** file.

```
--gene-matrix ARG
```

You must provide a single gene matrix that must be in [GCT](#) format.

```
zcat NovartisGeneAtlas2004.gct.gz | cut -f1-4 | head
```

```
#1.2
17581 79
Name Description Colorectal_Adenocarcinoma Whole_Blood
1 A1BG 115.5 209.5
2 A2M 85 328.5
9 NAT1 499 1578
10 NAT2 115 114
12 SERPINA3 419.5 387.5
13 AADAC 125 252.5
14 AAMP 2023 942.5
```

```
--condition ARG (Optional)
```

You may provide column names present in the **--gene-matrix** file, one per line. The matrix will be conditioned on these columns before the analysis is performed to help you identify secondary signals independent of these columns. Binary (0, 1) matrices will not be conditioned.

```
head conditions.txt
```

```
Whole_Blood
```

--gene-intervals ARG

You must provide gene intervals in BED format with a fourth column that contains the same gene identifiers as those present in the Name column of the **--gene-matrix** GCT file. Only the first four columns are used.

```
zcat NCBIgenes2013.bed.gz | head
```

```
chr1 10003485 10045555 64802 NMNAT1
chr1 100111430 100160096 54873 PALMD
chr1 100163795 100164756 100129320 HMGB3P10
chr1 100174205 100232185 391059 FRRS1
chr1 10027438 10027515 100847055 MIR5697
chr1 100308165 100308317 100270894 RPL39P9
chr1 100315632 100389578 178 AGL
chr1 100433941 100435837 730081 LOC730081
chr1 100435344 100492534 23443 SLC35A3
chr1 100503669 100548932 64645 HIAT1
```

--snp-intervals ARG

SNP linkage intervals must be specified in BED format and include a fourth column with the SNP identifiers. The linkage intervals assigned to the trait-associated SNPs you provide with **--snps** are taken from this file.

```
zcat TGP2011.bed.gz | head
```

```
chr1 0 254996 rs113759966
chr1 0 254996 rs114420996
chr1 0 254996 rs114608975
chr1 0 254996 rs115209712
chr1 0 254996 rs116400033
chr1 0 254996 rs116504101
chr1 0 254996 rs12184306
chr1 0 254996 rs12184307
chr1 0 254996 rs138808727
chr1 0 254996 rs139113303
```

--null-snps ARG

The null SNPs file must have one SNP identifier per line. Only the first column is used. The identifiers must be a subset of the identifiers in **--snp-intervals**.

```
zcat Lango2010.txt.gz | head
```

```
rs58108140 chr1 10583
```

```
rs180734498 chr1    13302
rs140337953 chr1    30923
rs141149254 chr1    54490
rs2462492   chr1    54676
rs10399749  chr1    55299
rs189727433 chr1    57952
rs149755937 chr1    59040
rs77573425  chr1    61989
rs116440577 chr1    63671
```

Output Files

The usage example shown above produces the following output files:

```
out/
  args.txt
  pvalues.txt
  null_pvalues.txt
  snp_genes.txt
  snp_pvalues.txt
```

args.txt

The command line arguments needed to reproduce the analysis.

```
cat args.txt

# SNPsea v1.0.2
--snps          LDL_Teslovich2010.txt
--gene-matrix   NovartisGeneAtlas2004.gct.gz
--gene-intervals NCBIGenes2013.bed.gz
--snp-intervals TGP2011.bed.gz
--null-snps     Lango2010.txt.gz
--out           out
--slop          250000
--threads       4
--null-snpsets  1000
--min-observations 50
--max-iterations 1000000
```

Repeat the analysis:

```
snpsea --args args.txt
```

pvalues.txt

The p-values representing enrichment of condition-specificity for the given SNPs.

```
head pvalues.txt | column -t
```

name	pvalue	nulls_observed	nulls_tested
ColorectalAdenocarcinoma	0.87	87	100
WholeBlood	0.00606299	77	12700
BM-CD33+Myeloid	0.054	81	1500
PB-CD14+Monocytes	0.316667	95	300
PB-BDCA4+Dendritic_Cells	0.164286	115	700
PB-CD56+NKCells	0.000104993	86	819100

null_pvalues.txt

If the argument for **--snps** is the name of a file, the p-values for null matched SNP sets. You can compare these null results to the results for your trait-associated SNPs.

If the argument for **--snps** is “randomN” where N is some integer, like “random20” the p-values for random unmatched SNP sets, each with N SNPs.

The fifth column is the replicate index. The number of replicates performed is specified with **--null-snpsets INT**.

```
head null_pvalues.txt | column -t
```

ColorectalAdenocarcinoma	0.056	84	1500	0
WholeBlood	0.236667	71	300	0
BM-CD33+Myeloid	0.55	55	100	0
PB-CD14+Monocytes	0.59	59	100	0
PB-BDCA4+Dendritic_Cells	0.59	59	100	0
PB-CD56+NKCells	0.71	71	100	0
PB-CD4+Tcells	0.383333	115	300	0
PB-CD8+Tcells	0.128571	90	700	0
PB-CD19+Bcells	0.168571	118	700	0
BM-CD105+Endothelial	0.386667	116	300	0

snp_genes.txt

Each SNP’s linkage interval and overlapping genes. If a SNP is not found in the reference file specified with **--snp-intervals**, then the name of the SNP will be listed and the other columns will contain NA.

```
head snp_genes.txt | column -t
```

chrom	start	end	name	n_genes	genes
-------	-------	-----	------	---------	-------

chr7	128560761	128773770	rs10488631	2	3663,23534
chr2	100637056	100895271	rs10865035	3	3899,150577,164832
chr11	118475098	118746223	rs10892279	2	1656,23187
NA	NA	NA	rs99999999	NA	NA
chr21	43817297	43851877	rs11203203	1	53347
chr1	117256697	117293763	rs11586238	4	914,965,3321,5738
chr1	161389417	161637888	rs12746613	5	2212,2213,2215,3310,9103
chr2	61068167	61382443	rs13031237	5	5194,5966,84542,339803,339804
chr3	58553160	58558769	rs13315591	1	11170

snp_pvalues.txt


Each SNP, condition, gene with greatest specificity to that sample, and p-value for the SNP-sample pair, adjusted for the number of genes overlapping the given SNP.

```
head snp_pvalues.txt | column -t
```

marker	column	gene	pvalue
rs10488631	ColorectalAdenocarcinoma	3663	0.270409
rs10488631	WholeBlood	3663	0.302693
rs10488631	BM-CD33+Myeloid	3663	0.0569547
rs10488631	PB-CD14+Monocytes	3663	0.0960891
rs10488631	PB-BDCA4+Dendritic_Cells	3663	0.240571
rs10488631	PB-CD56+NKCells	23534	0.58674
rs10488631	PB-CD4+Tcells	3663	0.683486
rs10488631	PB-CD8+Tcells	23534	0.634216
rs10488631	PB-CD19+Bcells	3663	0.261931

Output Visualizations


View enrichment of tissue-specific gene expression



Create this visualization with:

```
python bin/snpsea-barplot --out out
```


View the most specifically expressed gene for each SNP-tissue pair



Create this visualization with:

```
python bin/snpsea-heatmap --out out
```

View the type 1 error rate estimates for each tissue



Create this visualization with:

```
Rscript bin/snpsea-type1error out
```